

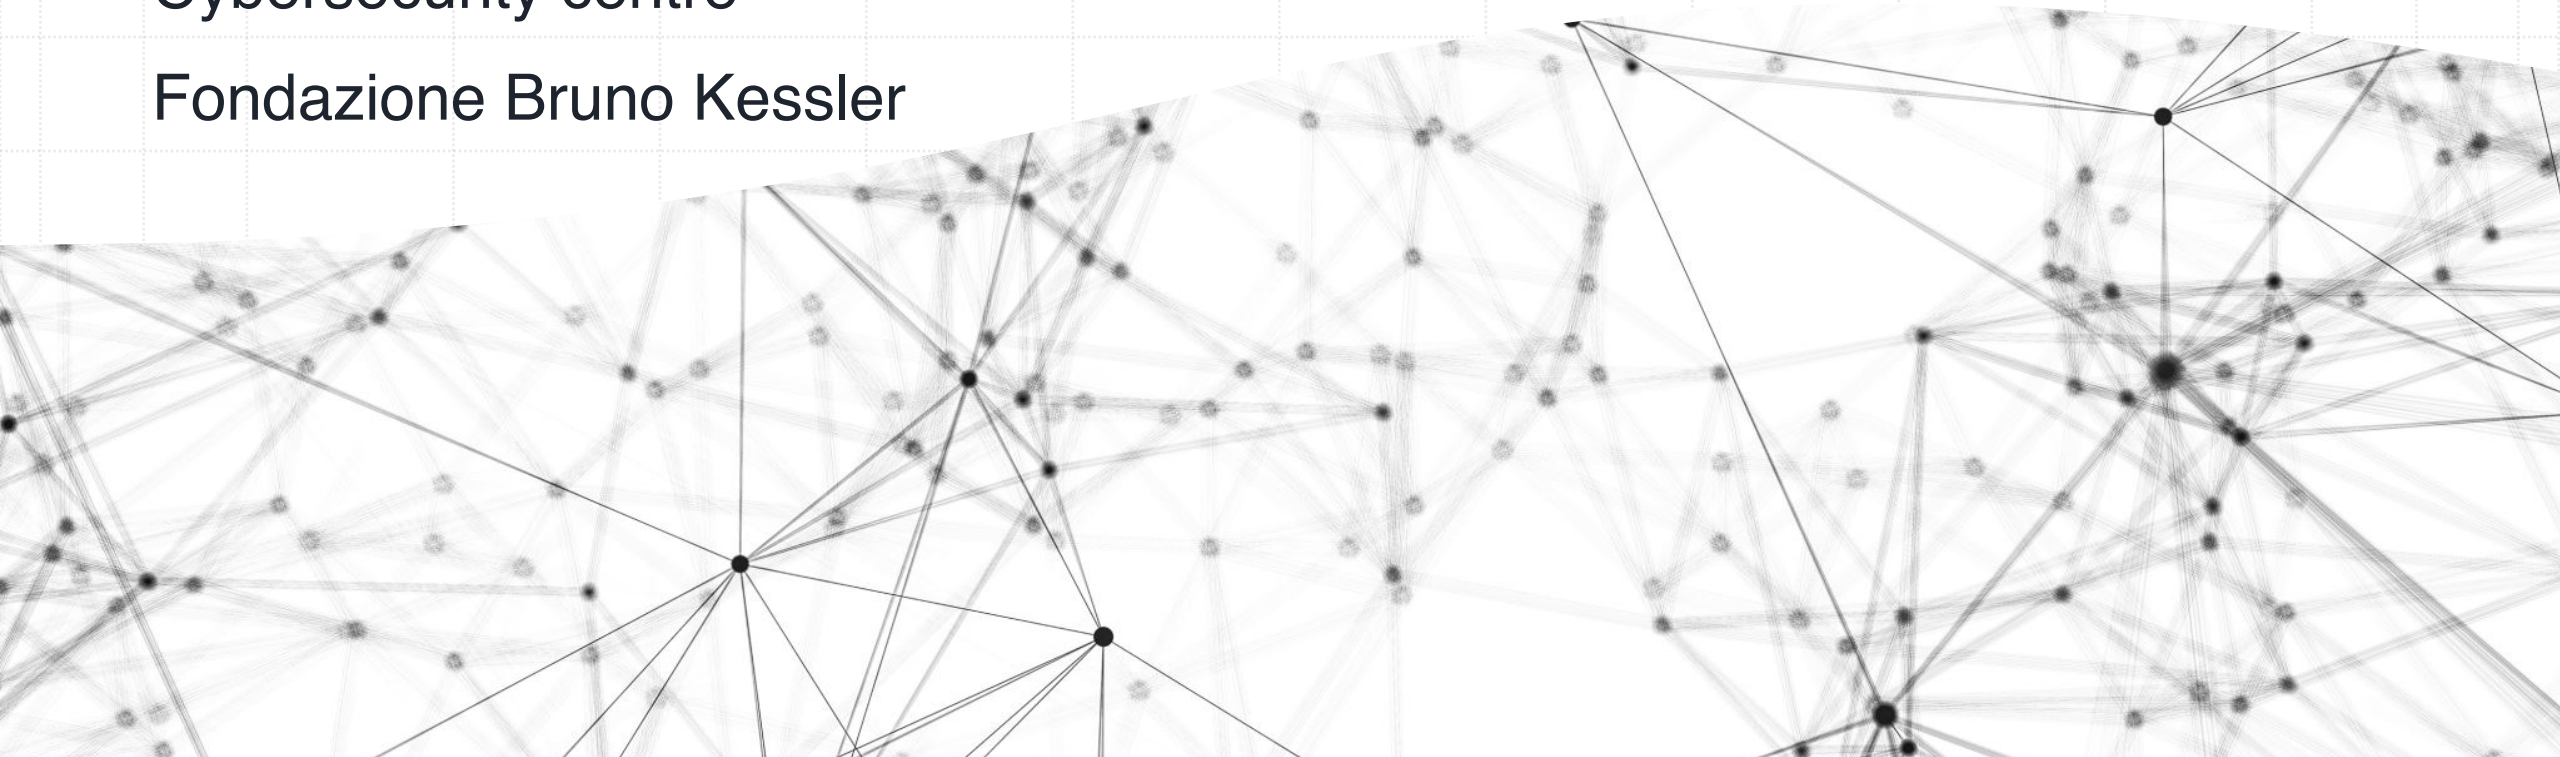


Traffic feature importance

Roberto Doriguzzi Corin

Cybersecurity centre

Fondazione Bruno Kessler





Outline

- Traffic feature engineering
 - Packet-level features
 - Flow-level features
- Techniques to expose relative feature importance and select relevant features



Feature engineering

Process that involves:

- **Feature selection** (selecting the most useful features)
- **Feature extraction**
 - extracting relevant features from the raw data
 - combining existing features to produce a more useful one
- **Feature processing**
 - Normalisation of the features into a common range (e.g., $[0,1]$)
 - Encoding categorical data into integer format (e.g., one-hot encoding)

Feature processing (1)

Normalisation

- Neural networks **converge much faster** with normalised features
- Normalisation (or *min-max scaling*) shifts and re-scales all the features to [0,1] range

$$x = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$$


$$w_j := w_j - \alpha \cdot \frac{\partial}{\partial w_j} J(W)$$

$$:= w_j - \alpha \cdot \frac{\partial}{\partial w_j} \frac{1}{2m} \sum_{i=1}^m (h_W(x^{(i)}) - y^{(i)})^2$$


$$:= w_j - \alpha \cdot \frac{1}{m} \sum_{i=1}^m (h_W(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)}$$

The presence of feature value x_j in the formula affects the step size of the gradient descent

Gradient descent without scaling



Gradient descent after scaling variables



Source: Stanford University



Feature processing (2)

One-hot-encoding

Categorical data are variables that contain label values rather than numeric values (e.g., 5 traffic classes: Benign, DDoS, Brute Force, Port Scan, SQL Injection).


Many machine learning algorithms cannot operate on text data directly. They require all input variables and output variables to be numeric.

Two steps:

- Integer encoding (e.g., Benign=0, DDoS=1, etc,)
- One-hot encoding (Benign=[1,0,0,0,0], DDoS=[0,1,0,0,0])



Feature extraction



Reminder: we are looking for malicious traffic

- Attack traffic can be generated intentionally or unintentionally.
- In the first case we might prefer to block the source host/network
- In the second case, we might want to block a specific application (malware installed inside our network)
- What is the best policy?
 - A fine-grained policy is more precise, but it requires more memory and CPU resource
 - A coarse-grained policy is easier to manage, but might prevent legitimate services to work
- Impact not only on the mitigation strategy
- But also on the way we process the traffic for intrusion detection

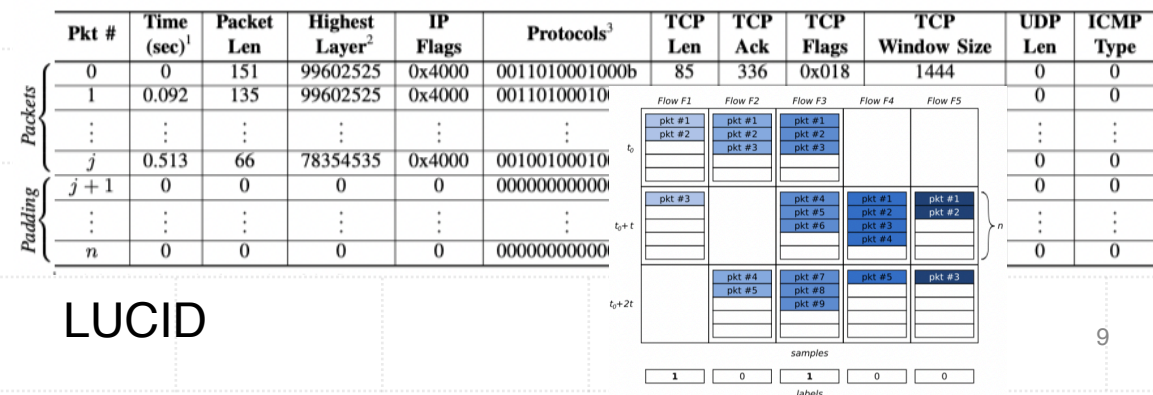
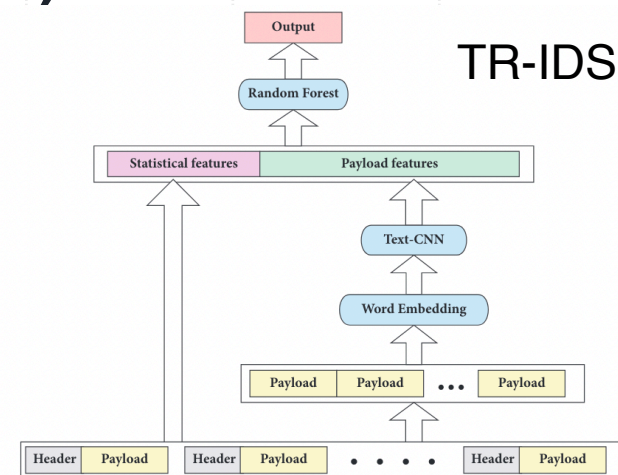


Flow-based processing

- The most common approach is that of collecting flow-specific features
 - A flow is usually defined by using a tuple like `(srcIP, dstIP, srcPort, dstPort, protocol)`
 - A flow can be represented using packet-based features or statistical features
 - Unidirectional or bi-directional?
 - Bidirectional gives a better view of the interaction between attacker and victim
 - However, it makes it harder to identify the source of the attack


Feature extraction (Packet-level features)

- Packet-level features:
 - Build a representation of a flow grouping representations of packets that belong to the same flow
 - Packet-level info such as header fields (sometimes payload as well)
- IP addresses and ports are usually avoided as features (too specific to the testbed where the data has been collected)



DeepDefense

Field	Field Example	Field Type
frame.encap_type	1	Boolean
frame.len	805	Numerical
frame.protocols	eth:ip:tcp:http:data:	Text

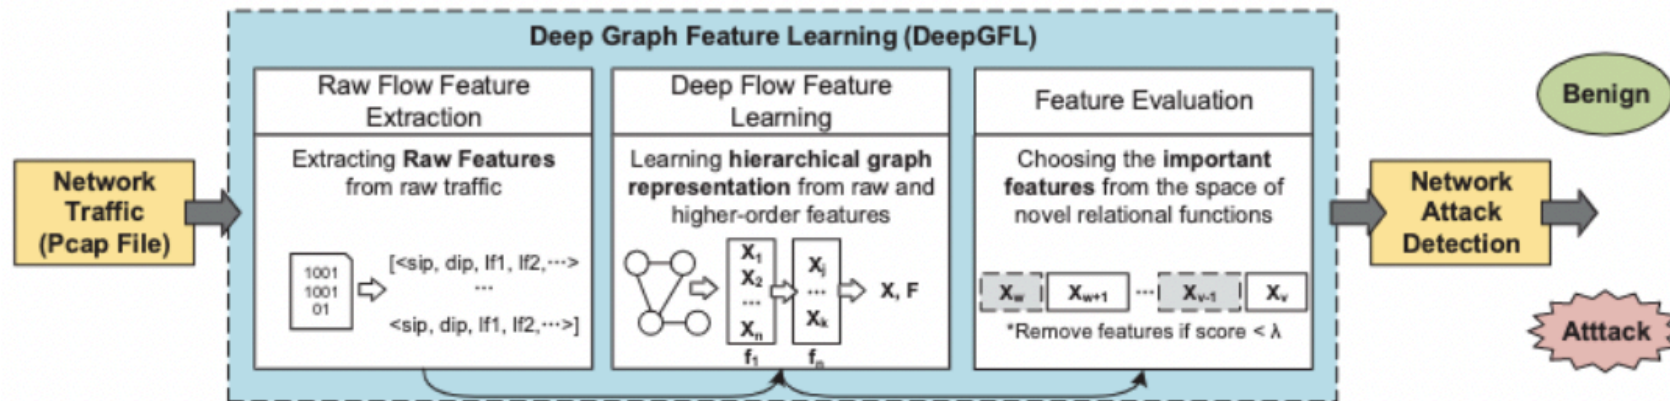


tcp.analysis.ack_rtt	0	Numerical
tcp.analysis.bytes_in_flight	1.460e+03	Numerical
tcp.analysis.duplicate_ack_num	1	Numerical
tcp.dstport	2090	Boolean
tcp.flags.urg	0	Boolean
tcp.len	751	Numerical
tcp.srcport	80	Boolean
tcp.window_size	12864	Numerical
udp.dstport	47666	Boolean
udp.length	97	Numerical
udp.srcport	47521	Boolean

Feature extraction (Statistical features)

- Flow-level features:
 - Flows are represented using statistical features
 - Many papers use pre-extracted features provided with the dataset

Feature	Description
protocol	Protocol of the flow
src_port	Source port
dst_port	Destination port
f(b)_urg_num	Number URG flags in the forward(backward) direction (0 for UDP)
f(b)_ack_num	Number ACK flags in the forward(backward) direction (0 for UDP)
f(b)_psh_num	Number PSH flags in the forward(backward) direction (0 for UDP)
f(b)_rst_num	Number RST flags in the forward(backward) direction (0 for UDP)
f(b)_syn_num	Number SYN flags in the forward(backward) direction (0 for UDP)
f(b)_fin_num	Number FIN flags in the forward(backward) direction (0 for UDP)
pkts_num	Total packets in the flow
bytes_num	Total bytes in the flow
f(b)_pkts_num	Total packets in the forward(backward) direction
f(b)_bytes_num	Total bytes in the forward(backward) direction
f(b)_len_min	Minimum length of packet in the forward(backward) direction
f(b)_len_max	Maximum length of packet in the forward(backward) direction
f(b)_len_mean	Mean length of packet in the forward(backward) direction
f(b)_len_std	Standard deviation length of packet in the forward(backward) direction
duration	Duration of the flow
pkts_psec	Number of packets per second
bytes_psec	Number of bytes per second
f(b)_pkts_psec	Number of forward(backward) packets per second
f(b)_bytes_psec	Number of forward(backward) bytes per second
f(b)_intv_min	Minimum time interval between two packets sent in the forward(backward) direction
f(b)_intv_max	Maximum time interval between two packets sent in the forward(backward) direction
f(b)_intv_mean	Mean time interval between two packets sent in the forward(backward) direction
f(b)_intv_std	Standard deviation time interval between two packets sent in the forward(backward) direction



DeepGFL. Feature extraction with CICFlowMeter



Remark

Usability of the traffic representations in real-world deployments

- In real-world application the traffic is collected within a **time-window** and the processed
- Tools like CICFlowMeter detect the beginning and the end of a flow for computing its statistical properties
 - However, we cannot wait the end of an attack to detect it. It would be too late!!!

Length of the time window?

- Depends on the application and on the available resources:
 - Long time-windows allow for **more detailed flow representations**, but require more memory and CPU resources
 - Short time-windows allow for **faster processing and reaction**, but less data is can be collected

Balancing between time-window, amount of incoming traffic and resources is often necessary

- **Dynamic time-windows, packet sampling** are techniques that can be used to solve the problem

Feature extraction with tshark and pyshark

- Tshark: <https://www.wireshark.org/docs/man-pages/tshark.html>
- Tshark lets you **capture** packet data from a live network, or **read** packets from a previously saved capture file, either **printing** a decoded form of those packets to the standard output or **writing** the packets to a file.
- Pyshark: <https://github.com/KimiNewt/pyshark>
- Python wrapper for tshark, allowing python packet parsing using wireshark dissectors (plugins to decode protocol encapsulations of packets).

Why tshark?

- Easy to debug: its output can be checked in Wireshark
- High-level packet features: e.g.: list of protocols.
- Can be used for live-testing (see LUCID code)
- Supports pcap file format -> compatible with tcpdump

Transmission Control Protocol, Src Port: 22607, Dst Port: 80, Seq: 0, Len: 0

```
Source Port: 22607
Destination Port: 80
[Stream index: 0]
[TCP Segment Len: 0]
Sequence number: 0 (relative sequence number)
Sequence number (raw): 1640917568
[Next sequence number: 1 (relative sequence number)]
```

> Acknowledgment number: 800978182

Acknowledgment number (raw): 800978182
0101 = Header Length: 20 bytes (5)

Flags: 0x002 (SYN)

```
000. .... = Reserved: Not set
...0 .... = Nonce: Not set
.... 0... = Congestion Window Reduced (CWR): Not set
.... 0... = ECN-Echo: Not set
.... ..0. = Urgent: Not set
.... ...0 = Acknowledgment: Not set
.... .... 0... = Push: Not set
.... .... .0.. = Reset: Not set
```

```
> .... ..1. = Syn: Set
```

```
.... 0 = Fin: Not set
```

[TCP Flags:S.]

Window size value: 512

```
01 protocols = {LayerFieldsContainer} eth:ethertype:ip:tcp
```

```
tcp = {Layer} TCP: \n\tSource Port: 22607\n\tDestination Port: 80\n\tStream index: 0\n\tSequence = {LayerFieldsContainer} Layer TCP: \n\tSource Port: 22607\n\tDestination Port: 80\n\tStream index: 0\n\tTCP Segment Len: 0
DATA_LAYER = (str) 'data'
ack = {LayerFieldsContainer} 800978182
ack_nonzero = {LayerFieldsContainer} The acknowledgment number field is nonzero while the ACK flag is not set
ack_raw = {LayerFieldsContainer} 800978182
checksum = {LayerFieldsContainer} 0x00004e73
checksum_status = {LayerFieldsContainer} 2
connection_syn = {LayerFieldsContainer} Connection establish request (SYN): server port 80
dstport = {LayerFieldsContainer} 80
field_names = (list: 37) ['srcport', 'dstport', 'port', 'stream', 'len', 'seq', 'seq_raw', 'nextseq', 'ack', '_ws_expert', 'ack_non
flags = {LayerFieldsContainer} 0x00000002
flags_ack = {LayerFieldsContainer} 0
flags_cwr = {LayerFieldsContainer} 0
flags_ecn = {LayerFieldsContainer} 0
flags_fin = {LayerFieldsContainer} 0
flags_ns = {LayerFieldsContainer} 0
flags_push = {LayerFieldsContainer} 0
flags_res = {LayerFieldsContainer} 0
flags_reset = {LayerFieldsContainer} 0
flags_str = {LayerFieldsContainer} \xc2\xb7\xc2\xb7\xc2\xb7\xc2\xb7\xc2\xb7\xc2\xb7\xc2\xb7\xc2\xb7\xc2\xb7\xc2\xb7\xc2\xb7
flags_syn = {LayerFieldsContainer} 1
flags_urg = {LayerFieldsContainer} 0
hdr_len = {LayerFieldsContainer} 20
layer_name = (str) 'tcp'
len = {LayerFieldsContainer} 0
nextseq = {LayerFieldsContainer} 1
port = {LayerFieldsContainer} 22607
raw_mode = (bool) False
seq = {LayerFieldsContainer} 0
seq_raw = {LayerFieldsContainer} 1640917568
srcport = {LayerFieldsContainer} 22607
stream = {LayerFieldsContainer} 0
time_delta = {LayerFieldsContainer} 0.000000000
time_relative = {LayerFieldsContainer} 0.000000000
urgent_pointer = {LayerFieldsContainer} 0
window_size = {LayerFieldsContainer} 512
window_size_value = {LayerFieldsContainer} 512
```


Feature extraction with Tshark

- Both packet-level and flow-level features can be extracted with tshark/pyshark
- Alternative tools are:
 - **CICFlowMeter** (<https://github.com/CanadianInstituteForCybersecurity/CICFlowMeter>), a Java-based tool for feature extraction (statistical features).
 - **Tcpdump and libpcap** (<https://www.tcpdump.org/>), a command line tool and a C/C++ library for network traffic capture and editing (e.g., split).
 - Other tools for pcap file manipulation are
 - **tcprewrite**: packet editing
 - **tcpreplay**: replay a pre-recorded pcap file through a network interface



Feature importance

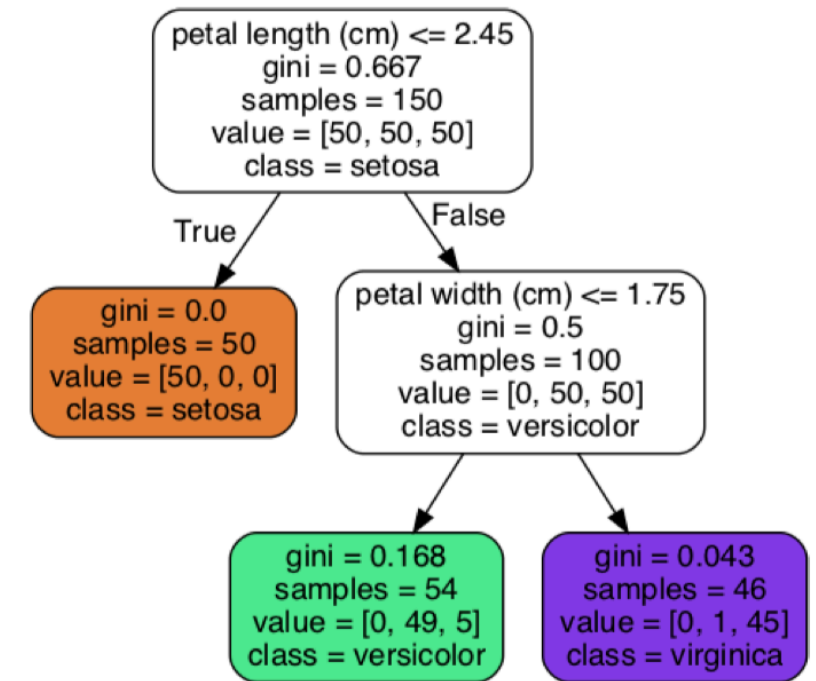
Decision trees(1)

- Decision trees are simple but powerful ML algorithms for classification and regression tasks
- Their decisions are easy to interpret:
 - samples** counts how many samples a node applies to
 - value** counts how many training samples of each class the node applies to
 - gini** is a measure of impurity (gini=0 means pure) and is computed as $G_i = 1 - \sum_{k=1}^n p_{i,k}^2$

Where $p_{i,k}$ is the ratio of class k samples among the training instances of the i^{th} node.

Example: the Gini impurity of the depth-2 right node is equal to:

$$1 - (0/46)^2 + (1/46)^2 + (45/46)^2 \approx 0.043$$



Iris Decision Tree (source: Géron, A. (2022). *Hands-on machine learning with Scikit-Learn, Keras, and TensorFlow*. " O'Reilly Media, Inc.")

Decision trees (2)

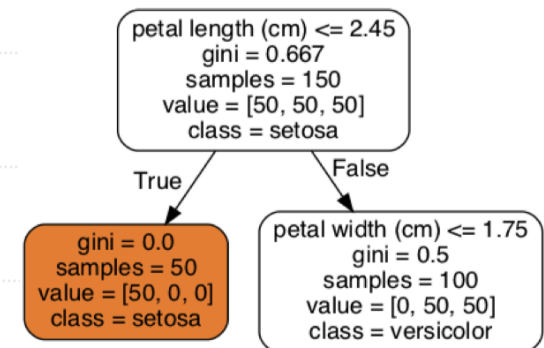
- Scikit-Learn uses the CART algorithm to train DTs
- CART splits the set of each node using a **single feature k and a threshold t_k**
- the algorithm chooses **the pair (k, t_k) that produces the purest subsets** (lowest gini)

CART cost function for classification:

$$J(k, t_k) = \frac{m_{\text{left}}}{m} \cdot G_{\text{left}} + \frac{m_{\text{right}}}{m} \cdot G_{\text{right}}$$

Where: $G_{\text{left/right}}$ measures the impurity of the left/right subset

$m_{\text{left/right}}$ is the number of instances in the left/right subset




Decision trees (3)

- CART recursively splits each subset until the max depth is reached or the impurity cannot be reduced with further splits
- For regression tasks, CART splits the training set in a way to minimise the MSE
- Decision trees are easy to interpret (often called *white box* models)

Limitations of DTs:

- Sensitive to training set rotation
- Sensitive to small variations of the training set (adding/removing training samples might results in totally different splits)



Source: <https://towardsdatascience.com/>



Random Forests can limit DT instability by averaging the prediction of many DTs

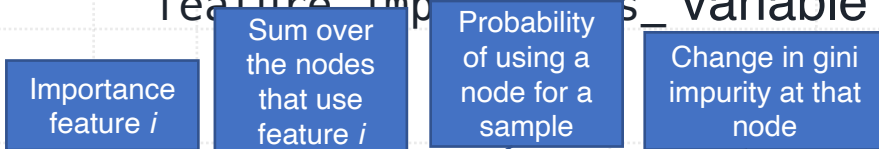


Random Forests

- A Random Forest is an ensemble of DTs generally trained via **bagging method** (sampling with replacement)
- When splitting a node, RF searches for the best feature among a **random subset of features** (instead of all to reduce the variance)
- Extremely Randomized Trees (Extra-Trees) use **random thresholds** when splitting a node. This reduces the training time (finding the best threshold is time-consuming)
- Classification can be done by assigning the class that obtains the majority of votes (**hard voting**), or assigning the class that obtains the highest average probability (**soft voting**)
- the **scikit-learn implementation** combines classifiers by averaging their probabilistic prediction (**soft voting**), instead of letting each classifier vote for a single class

Feature selection with Random Forest

- Scikit-learn measures a feature's importance by looking at how much the tree nodes that use that feature reduce impurity
 - On average across all the tree in the forest
- The values are accessible using the `feature_importances_` variable



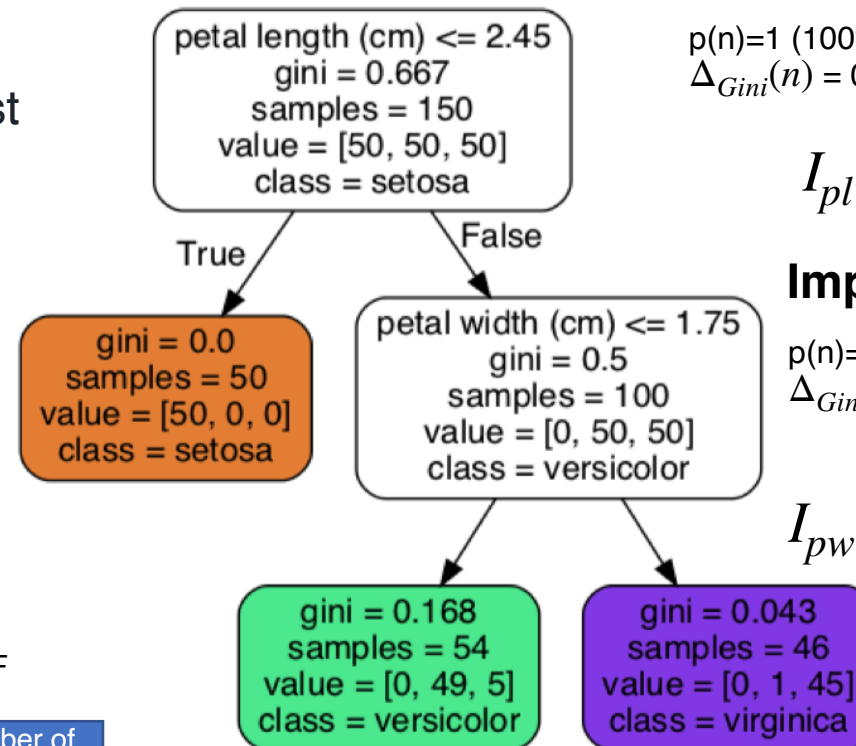
$$I_i(T) = \sum_{n \in T, i_n = i} p(n) \Delta_{Gini}(n)$$

Feature importance with one DT
(n =node, i =feature)

$$I_i = \frac{1}{|B|} \sum_{T \in B} I_i(T)$$

Feature importance with RF

Number of trees



Importance of petal length

$$p(n)=1 \text{ (100\% of samples)}$$

$$\Delta_{Gini}(n) = 0.667 - (0.334 \cdot 0 + 0.667 \cdot 0.5) = 0.33$$

$$I_{pl} = 1 \cdot 0.33 = 0.33$$

Importance of petal width

$$p(n)=0.667$$

$$\Delta_{Gini}(n) = 0.5 - (0.54 \cdot 0.168 + 0.46 \cdot 0.043)$$

$$= (0.09 + 0.02) = 0.39$$

$$I_{pw} = 0.667 \cdot 0.39 = 0.26$$



Laboratory: feature selection with RF

- Your problem is to **decide which are the relevant features** for your DDoS detection
- The representation of a flow is in an array-like format, where each row contains a representation of a packet, while each column is a packet-header feature
- A hidden feature is the length of the flow, thus the number of non-zero rows
- This laboratory consists of
 1. transforming the array-like representation of flows into 1-dimensional vectors, with one position representing the length of the flow
 2. Exploring different options for sample's dimensionality reduction (consider the meaning of a feature)
 3. Computing the feature importance with RF



Laboratory: Analysis of model's behaviour

- Once the model is trained, you might want to **understand how your IDS works**
- A basic way to do that is checking which **features are more important** for the classification
- One option is to **set to zero one feature at a time**, a measure the change in model's performance (compare to using all the features)
- If done for all the features, one can understand which ones are more important for the classification
- In the laboratory, you will **re-use the 1D representation** of the flows on the test set to understand which features are more important for a pre-trained model.